

Grammarinator: A Grammar-Based Open Source Fuzzer

Renáta Hodován
University of Szeged
Department of Software Engineering
Szeged, Hungary
hodovan@inf.u-szeged.hu

Ákos Kiss
University of Szeged
Department of Software Engineering
Szeged, Hungary
akiss@inf.u-szeged.hu

Tibor Gyimóthy
University of Szeged
MTA-SZTE Research Group on
Artificial Intelligence
Szeged, Hungary
gyimothy@inf.u-szeged.hu

ABSTRACT

Fuzzing, or random testing, is an increasingly popular testing technique. The power of the approach lies in its ability to generate a large number of useful test cases without consuming expensive manpower. Furthermore, because of the randomness, it can often produce unusual cases that would be beyond the awareness of a human tester. In this paper, we present Grammarinator, a general purpose test generator tool that is able to utilize existing parser grammars as models. Since the model can act both as a parser and as a generator, the tool can provide the capabilities of both generation and mutation-based fuzzers. The presented tool is actively used to test various JavaScript engines and has found more than 100 unique issues.

CCS CONCEPTS

• **Software and its engineering** → **Software testing and debugging**;

KEYWORDS

fuzzing, random testing, grammars, security

ACM Reference Format:

Renáta Hodován, Ákos Kiss, and Tibor Gyimóthy. 2018. Grammarinator: A Grammar-Based Open Source Fuzzer. In *Proceedings of the 9th ACM SIGSOFT International Workshop on Automating TEST Case Design, Selection, and Evaluation (A-TEST '18)*, November 5, 2018, Lake Buena Vista, FL, USA. ACM, New York, NY, USA, 4 pages. <https://doi.org/10.1145/3278186.3278193>

1 INTRODUCTION

Fuzzing [6], or random testing, is a popular technique as it promises the creation of a large number of test cases with limited effort. Moreover, as a result of randomness, it is often capable of generating extreme test cases that are easily overlooked by a human test engineer. Of course, there is a limit to the application area of such ‘blindly’ generated test cases as it is difficult to decide in the general case whether the program that consumed the random input

executed in a semantically correct way. However, some important aspects of the execution can still be monitored, such as whether the program accessed resources (e.g., code, stack, heap memory) according to its permissions or adhered to its design contracts (e.g., assertions). This makes the technique highly useful for security testing [8].

The first fuzzers were truly random, simply feeding random byte sequences to their system-under-test (SUT). Although such simple fuzzers can reveal some surprising faults, they most often cannot scratch anything but the surface of SUTs that have a complex input structure. The test cases generated by such format-unaware fuzzers are usually caught and discarded by the input parser, so deeper code paths do not get exercised. Model-based fuzzers aim at reaching these deeper parts of the SUT by generating test cases from some model of the input format, trying to ensure that the high-level structures that get checked by the SUT’s parser are valid while the ‘real’ content gets randomized. Model-based fuzzers are usually sorted into two main categories based on how they build their model and generate test cases from it: purely generation-based fuzzers expect an explicit representation of the input format, usually written by hand, while purely mutation-based approaches need a suite of existing test cases, which act as an implicit representation of the input format and can be mutated or recombined to yield new inputs. Both approaches have their advantages and disadvantages. The building of a model of a complex input format for a purely generation-based fuzzer may be a time-consuming, tedious task: e.g., generating random programs as test inputs for compilers or interpreters needs the modelling of syntax, symbol tables, type information, etc. On the other hand, mutation-based fuzzers require a vast number of existing test cases to work with (and they may still need additional information on where and how to mutate or recombine them).

In this paper, we present a tool named Grammarinator that aims at providing the capabilities of both generation and mutation-based fuzzers with the help of grammars. Grammarinator can turn existing parser grammars into generation models while using the same grammars to analyse and evolve existing test sets. The tool is being actively used for the testing of JavaScript engines and has found more than 100 unique defects.

The rest of the paper is organised as follows: Section 2 describes the challenges of grammar-based random test generation and how Grammarinator solves them. Section 3 reports on the results of Grammarinator used in practice. Section 4 highlights some related work from the academia and from the industry. And finally, Section 5 summarises the paper.

Permission to make digital or hard copies of all or part of this work for personal or classroom use is granted without fee provided that copies are not made or distributed for profit or commercial advantage and that copies bear this notice and the full citation on the first page. Copyrights for components of this work owned by others than the author(s) must be honored. Abstracting with credit is permitted. To copy otherwise, or republish, to post on servers or to redistribute to lists, requires prior specific permission and/or a fee. Request permissions from permissions@acm.org.

A-TEST '18, November 5, 2018, Lake Buena Vista, FL, USA

© 2018 Copyright held by the owner/author(s). Publication rights licensed to ACM.

ACM ISBN 978-1-4503-6053-1/18/11...\$15.00

<https://doi.org/10.1145/3278186.3278193>




Figure 1: Overview of Grammarinator's approach to generate test cases.

2 GRAMMARINATOR

The primary goal of Grammarinator¹ is to be a generic grammar-based random test generator that is useful out-of-the-box for simple input formats while remaining customizable to support formats with complex requirements. Both simplicity and extensibility are achieved by adapting parsing concepts to generation purposes.

Generation from grammars. Traditional parser generator systems take a grammar and create two software components from it: a lexer to tokenize an input character stream and a parser to build an abstract syntax tree (AST) from the tokens according to the grammar rules. Then, these two components can work together to process any (correct) input. Grammarinator is designed to *mirror* these concepts: given a grammar, it creates two software components that are capable of generating input. To highlight the mirroring of the traditional concepts, the token generator part is called *unlexer* while the syntactic structures are generated by an *unparser*. Their combined result is an AST, which can be serialized into an actual test case.

The existence of grammars is key to the usefulness of this approach. To facilitate putting in practice, Grammarinator uses the widespread ANTLR v4 grammar syntax to model the input format. The central grammar repository of the ANTLR v4 tool² and the more than 150 open-source grammar definitions therein save a lot of modelling effort and open up the possibility of fuzz testing of a large number of SUTs.

Figure 1 gives a high-level overview of the workflow of Grammarinator: the top-left part depicts the (one-off) creation of the generator components and the right-hand side of the diagram visualizes the generation of ASTs and their serialization to test cases (some yet-undiscussed parts of the chart will be detailed later).

Listing 1 shows an example grammar of simple additive expressions and variable assignments (which can also be seen as an extremely restricted subset of JavaScript), while Listings 3 and 4 give excerpts from the Grammarinator-generated Python 3 sources of the unparser and unlexer. The code snippets show how the parser and lexer rules of the grammar are mapped to methods (e.g., defined at Lst. 3 line 17 and Lst. 4 line 36, used at Lst. 3 lines 26 and 77), how

Listing 1: Example grammar

```
1 grammar Example;
2 @parser::header{from random import sample}
3 program : {self.symtab = set(); (vardef=ID '=' expression '; '
4           {self.symtab.add(str($vardef))}); } ;
5 expression : value (('+'|'-') value)* | (('+'|'-')? value ;
6 value : '(' expression ')'
7         | {len(self.symtab) != 0}? ID {current.last_child =
8           UnlexerRule(src=sample(self.symtab, 1)[0])}
9         | NUM ;
10 ID : [a-z]+ ;
11 NUM : {0..1}? '0' | [1-9] [0-9]* ;
12 WS : [ \t\r\n]+ -> skip;
```

Listing 2: Test cases generated from grammar

```
1 m=+(o);uy=((((-+(0))))+mgbwf-3/(3+(-0)+(sk)));y=+pr;q=+0;
2 p=+(3-3);c=p;i=+c;rwke=(7-((i+0-p)-934)-i);aj=5+52-(0+p);z=rwke;t=+(+2);
```

the AST is built (e.g., child node created and added to its parent at Lst. 3 line 25), how rule alternatives are handled (e.g., selected at Lst. 4 lines 38, 40, and 42), and how quantifiers are dealt with (e.g., the Kleene plus at Lst. 3 line 22).

Customization. The first line of Listing 2 shows an output sample generated from the simplest version of the example grammar (the part of Listing 1 typeset in solid black). Assuming that the input format has requirements that are typical among programming languages but cannot be expressed using context-free grammars, like the forbidden use of undefined identifiers, it is clear that the generated test case is not perfect (i.e., its structure is seemingly correct but it will be probably discarded early on by a hypothetical SUT). Thus, to give more control over the generated content, Grammarinator also borrows the concepts of rule element labels, actions, and predicates from the ANTLR v4 syntax. Labels can capture generated subtrees for later reuse, actions allow the injection of arbitrary code into the generators, while predicates can guard alternatives during generation (also with the help of injected expressions).

The grey additions to the grammar in Listing 1 and their counterparts in the unparser and unlexer in Listings 3 and 4 exemplify how these constructs can be used to avoid undefined identifier references in generated test cases. I.e., at the beginning of the generation of a

¹<https://github.com/renatahodovan/grammarinator> or pip install grammarinator

²<https://github.com/antlr/grammars-v4>

Listing 3: Excerpt from unparser

```

4  from grammarinator.runtime import *
7  from random import sample
9  class ExampleUnparser(Grammarinator):
10     def __init__(self, unlexer):
11         super(ExampleUnparser, self).__init__()
12         self.unlexer = unlexer
13
14     @depthcontrol
15     def program(self):
16         local_ctx = dict()
17         current = self.create_node(UnparserRule(name='program'))
18         self.symtab = set()
19         if self.unlexer.max_depth >= 0:
20             for _ in self.one_or_more():
21                 current += self.unlexer.ID()
22                 local_ctx['vardef'] = current.last_child
23                 current += self.create_node(UnlexerRule(src='='))
24                 current += self.expression()
25                 current += self.create_node(UnlexerRule(src=';'))
26                 self.symtab.add(str(local_ctx['vardef']))
27
28         return current
29
30     program.min_depth = 3
31
32     @depthcontrol
33     def value(self):
34         current = self.create_node(UnparserRule(name='value'))
35         choice = self.choice([0 if [3, 1, 1][i] > self.unlexer.max_depth else w *
36                               self.unlexer.weights.get(('alt_25', i), 1) for i, w in enumerate([1,
37                                                  len(self.symtab) != 0, 1])])
38         self.unlexer.weights[('alt_25', choice)] = self.unlexer.weights.get(('alt_25',
39                                     choice), 1) * self.unlexer.cooldown
40
41         if choice == 0:
42             current += self.create_node(UnlexerRule(src='('))
43             current += self.expression()
44             current += self.create_node(UnlexerRule(src=')'))
45         elif choice == 1:
46             current += self.unlexer.ID()
47             current.last_child = UnlexerRule(src=sample(self.symtab, 1)[0])
48         elif choice == 2:
49             current += self.unlexer.NUM()
50
51         return current
52
53     value.min_depth = 1

```

Listing 4: Excerpt from unlexer

```

4  from grammarinator.runtime import *
5  charset_1 = list(chain(range(49, 58)))
6  charset_2 = list(chain(range(48, 58)))
7  class ExampleUnlexer(Grammarinator):
8     def __init__(self, *, max_depth=float('inf'), weights=None, cooldown=1.0):
9         super(ExampleUnlexer, self).__init__()
10         self.unlexer = self
11         self.max_depth = max_depth
12         self.weights = weights or dict()
13         self.cooldown = cooldown
14
15     @depthcontrol
16     def NUM(self):
17         current = self.create_node(UnlexerRule(name='NUM'))
18         choice = self.choice([0 if [0, 0][i] > self.unlexer.max_depth else w *
19                               self.unlexer.weights.get(('alt_1', i), 1) for i, w in enumerate([0, 1])])
20         self.unlexer.weights[('alt_1', choice)] = self.unlexer.weights.get(('alt_1',
21                                     choice), 1) * self.unlexer.cooldown
22
23         if choice == 0:
24             current += self.create_node(UnlexerRule(src='0'))
25         elif choice == 1:
26             current += self.create_node(UnlexerRule(src=self.char_from_list(
27                                     charset_1)))
28             if self.unlexer.max_depth >= 0:
29                 for _ in self.zero_or_more():
30                     current += self.create_node(UnlexerRule(src=self.char_from_list(charset_2)))
31
32         return current
33
34     NUM.min_depth = 0

```

‘program’, a simple symbol table is initialized (Lst. 1 line 3, Lst. 3 line 20), and whenever a new assignment is generated, the newly created token node is saved into a label (Lst. 1 line 3, Lst. 3 line 24) and added to the symbol table (Lst. 1 line 3, Lst. 3 line 28). If a ‘value’ is to be generated, a predicate is consulted whether the symbol table contains any identifiers that can appear on the right-hand side of an assignment (Lst. 1 line 6, Lst. 3 line 67). Finally, if the guard enables that alternative and the unparser chooses it, an identifier is picked randomly from the symbol table (Lst. 1 line 6, Lst. 3 line 75). The second line of Listing 2 gives a test case generated with the help of these additions, not suffering from dangling references.

There may be challenges that may not be solved (elegantly) with code fragments injected into the generators, or the extension of the grammar with such constructs is to be avoided for maintenance reasons. For these cases, Grammarinator adds a transformation step between AST generation and serialization, as shown in Figure 1, where totally arbitrary user-defined tree rewrites can be applied. Perhaps the simplest use case of this feature is the insertion of whitespace tokens in the output (which may be a stylistic improvement only, as for our example format, but may also be a strict necessity in more complex cases). Moreover, tree transformers could also be used to give an alternative solution for the identifier definition-reference agreement problem.

Controlling generation. A general challenge of random test generation is to avoid the creation of an excessively big content, also known as the bloating problem, which can easily happen with recursive models like grammars. A typical approach is to define an artificial size threshold and simply discard the test cases exceeding this limit. However, it is not worth wasting computational power on thrown-away test cases. Thus, Grammarinator takes on the approach of pre-calculating the minimum height of subtrees that the unlexer and unparser may generate for the rules of the grammar. This information is used during test case generation to guide alternative selection: only such alternatives are considered that can guarantee not to generate AST nodes deeper than a given threshold.

Listings 3 and 4 show examples of the precomputed minimum tree size information (e.g., Lst. 3 line 31, Lst. 4 line 49), of the book-keeping of the tree size during generation (e.g., Lst. 4 line 17, Lst. 3 line 16), and how this is used during alternative selection (e.g., Lst. 3 lines 21 and 67).

Another challenge of fuzzing with grammars is to cover all rules and alternatives – even those that can only be reached after several consecutive decisions – as balanced as possible without systematically generating all variants (which is impossible in finite time in the general case). Therefore, Grammarinator uses weighted random choice whenever generation has to choose between alternatives. By default, all alternatives have the same weight associated, but every time one of them is selected, its weight can be scaled (cooled) down by a user-defined factor. This approach helps guiding the generation toward the less-visited parts of the grammar.

Listings 3 and 4 demonstrate how this technique is manifested in the generators: how weights and the cooldown factor are initialized (Lst. 4 lines 18–19), how weights affect random decisions (e.g., Lst. 3 line 67), and how they are updated (e.g., Lst. 3 line 68).

Weighted choices can also be controlled manually. Predicates may not only guard alternatives with strictly boolean decisions but may also define weights to put more or less emphasis on some

branches. The listings show an example to guide numeric literals away from zero by explicitly giving a small weight in a predicate (Lst. 1 line 9, Lst. 4 line 38).

Reusing existing test cases. The advantage of grammars as fuzzer models is that they can be used to parse existing – and typically both syntactically and semantically correct – test cases. This enables Grammarinator to build a pool of ASTs and use them to boost purely generation-based test case creation with evolutionary techniques. One possibility is to perform a random recombination on trees from the pool to produce new test cases. But the generation and recombination-based approaches may also be fused together by taking trees from the pool and mutating them by replacing some parts with randomly generated new subtrees. The bottom-left part of Figure 1 depicts the place of these techniques in the workflow of Grammarinator.

3 RESULTS

Grammarinator was used to generate test inputs in ECMAScript format for the JerryScript³ engine, using the grammar available from the ANTLR v4 repository and applying the techniques described above. Depth control was used to keep a check on the size of the generated test cases, edge weight cooling was ensuring the diversity of the inputs, predicates and actions were injected to help meet the symbolic (e.g., no reference before definition) and structural (e.g., no loop control outside loops) requirements of the language, and transformers were adding whitespaces to the output. The conformance and regression test suites of the engine were used to build a population of trees for recombination and mutation. Guided by the Fuzzinator⁴ framework – which manages the passing of test cases from the fuzzer to the engine, the detection of erroneous behaviour, the classification of issues as unique or duplicate, and the minimization of failure-inducing inputs [1] –, Grammarinator has triggered 89 unique issues confirmed as valid by the JerryScript project team.

The tool was also used to fuzz the IoT.js⁵ project, a modular event-based JavaScript platform powered by the JerryScript engine. Although format (i.e., JavaScript) processing is delegated to JerryScript and IoT.js does not perform a lot of processing tasks itself, Grammarinator has revealed 15 unique issues in the platform, too.

All found issues have been reported in the issue trackers of the projects and are also listed in the wiki of Grammarinator⁶.

4 RELATED WORK

The idea of using grammars as test generation models has several decades of history. As early as in the '70s, Purdom [7] experimented with testing parser programs with test cases generated from context-free grammars. Later in the '90s, Maurer [5] used context-free grammars to generate test cases for VLSI circuits.

One of the closest approach to our work is LangFuzz [2] by Holler et al. It also parses existing regression tests with ANTLR grammars and caches all the subtrees into a fragment pool to be used for recombination during generation. It also aims to use the

grammars as a generator model, however, it supports only a subset of the ANTLR feature set. More complex grammar constructs like quantifiers or sub-alternatives are not supported and require manual grammar transformations. Moreover, the prototype tool of this research is not available as open source.

Kifetew et al. [3] use stochastic context-free grammars in BNF-format as test generation model. The probabilities of alternatives are gathered automatically from existing test cases to promote valid sentence structures and to control recursion depth. Additionally, they support evolutionary operators. In a follow-up work [4], they also discuss manual annotations. Unfortunately, their prototype tool was not available for public use.

A commercially utilized model-based fuzzing approach is that of Peach⁷. It has its own XML-based model format used only by the Peach Fuzzer. Additionally, because of its commercial focus, its open community edition has not been maintained since 2014.

5 SUMMARY

In this paper, we presented an open source tool, named Grammarinator, which is able to generate syntactically correct test cases from grammars. But, to be able to deal with not only the simplest cases, it can do more. It can limit the depth of generated structures and direct the generation toward less visited paths. By inlining code snippets into the grammar, it allows to define complex actions and decisions that a context-free grammar would not be able to describe. And with the help of tree transformers, it can manipulate test cases arbitrarily. Finally, Grammarinator can also exploit the fact that the same grammar that can generate new tests can also be used to parse existing test suites, and then create new content resulting from their recombination or mutation.

The tool has proven its usefulness in the hardening of real-life projects by revealing more than 100 valid unique issues.

ACKNOWLEDGMENTS

This research was partially supported by the Ministry of Finance of Hungary under EU-supported Hungarian national grant GINOP-2.3.2-15-2016-00037.

REFERENCES

- [1] Renáta Hodován and Ákos Kiss. 2018. Fuzzinator: An Open-Source Modular Random Testing Framework. In *Proceedings of the 11th IEEE International Conference on Software Testing, Verification and Validation (ICST 2018)*. IEEE, 416–421.
- [2] Christian Holler, Kim Herzig, and Andreas Zeller. 2012. Fuzzing with Code Fragments. In *Proceedings of the 21st USENIX Security Symposium (USENIX Security '12)*. USENIX Association, 445–458.
- [3] Fitsum Meshesha Kifetew, Roberto Tiella, and Paolo Tonella. 2014. Combining Stochastic Grammars and Genetic Programming for Coverage Testing at the System Level. In *Search-Based Software Engineering – 6th International Symposium, SSSE 2014, Proceedings*. Springer, 138–152.
- [4] Fitsum Meshesha Kifetew, Roberto Tiella, and Paolo Tonella. 2017. Generating valid grammar-based test inputs by means of genetic programming and annotated grammars. *Empirical Software Engineering* 22, 2 (2017), 928–961.
- [5] Peter M. Maurer. 1990. Generating Test Data with Enhanced Context-Free Grammars. *IEEE Software* 7, 4 (1990), 50–55.
- [6] Barton Miller. 2008. Foreword for Fuzz Testing Book. <http://pages.cs.wisc.edu/~bart/fuzz/Foreword1.html>.
- [7] Paul Purdom. 1972. A sentence generator for testing parsers. *BIT Numerical Mathematics* 12, 3 (1972), 366–375.
- [8] Ari Takanen, Jared DeMott, Charlie Miller, and Atte Kettunen. 2018. *Fuzzing for Software Security Testing and Quality Assurance* (2nd ed.). Artech House.

³<https://github.com/jerryscript-project/jerryscript>

⁴<https://github.com/renatahodovan/fuzzinator>

⁵<https://github.com/Samsung/iotjs>

⁶<https://github.com/renatahodovan/grammarinator/wiki>

⁷<https://www.peach.tech/>